

Constrained Optimizaton Modelling on Resource Allocation for Shared Hosting Platform

Teng Yu

Fall 2015

Abstract

The abstract is a very brief summary of the report's contents. It should be about half page long. Somebody unfamiliar with your project should have a good idea of what it is about having read the abstract alone and will know whether it will be of interest to them.

This work intend to improve the performance of resource allocating process on shared hosting cloud platform. We research the relations between different resource allocation requirements and compare with the structure inside the cluster. Then involving order theory and ranking functions to model the allocation and constraints more efficient instead of framing the process as a special Linear Programming formulation[Sti12]. Finally, we say this approach can be used to extend the work on dynamic workloads.

Acknowledgements

It is usual to thank those individuals who have provided particularly useful assistance, technical or otherwise, during your project. Your supervisor will obviously be pleased to be acknowledged as they will have invested quite a lot of time overseeing your progress.

Sincere thanks to Dr.Mark Stillwell for his countless helpful comments and advise during this project. Thanks Prof.Alexander Wolf for his supervision. Thanks my personal tutor Dr.Sadri Fariba and my friends in Imperial College for their helps.

Contents

1	Introduction	3
2	Background and Related Work	3
3	Body of report	4
3.1	Original Idea	4
4	Evaluation	5
5	Conclusions and Future work	5

This should list the main chapters and (sub) sections of your report. Choose self-explanatory chapter and section titles. If possible you should include page numbers indicating where each chapter/section begins. Try to avoid too many levels of subheading. Try if possible to stick to sections and subsections; subsubsections are usually avoidable.

1 Introduction

This is one of the most important components of the report. It should begin with a clear statement of what the project is about so that the nature and scope of the project can be understood by the reader. It should summarise everything you set out to achieve, provide a clear summary of the project's background and relevance to other work, and give pointers to the remaining sections of the report that contain the bulk of the technical material.

2 Background and Related Work

The background and related work section of the report should set the project into context by relating it to existing published work that you read at the start of the project when your approach and methods were being considered. There are usually many ways of approaching a given problem, and you should not just pick one at random. Describe and evaluate as many alternative approaches as possible. The published work may be in the form of research papers, articles, text books, technical manuals, or even existing software or hardware of which you have had hands-on experience. Do not be afraid to acknowledge the sources of your inspiration; you are expected to have seen and thought about other people's ideas, so your contribution largely will be putting them into practice in some other context.

3 Body of report

The central part of the report typically consists of three of four chapters detailing the technical work undertaken during the project. The structure of these chapters is highly project dependent. Usually they reflect the chronological development of the project, e.g., design, implementation, experimentation, and optimisation, although this is not always the best approach. However you choose to structure this part of the report, you should make it clear how you arrived at your chosen approach in preference to the other alternatives documented in the background. For implementation projects you should describe and justify the design of your system at some high level, for example by using any of the design methods taught during the first- and second-term courses, and should document any interesting problems with, or features of, your implementation. Integration and testing are also important to describe. Your supervisor will advise you on the most suitable structure for these middle sections.

3.1 Original Idea

Consider to handle with resource that have difficult-to-represent capacities. User's request may be special on different processing element (heterogeneous) or say the relation between them. Then we find the capacity of a cluster is the set of all allocation request that it can service. Try to give a relation between resource allocation requests (RAR):

Definition 1 *Given the relation \preceq between RAR:*

$$\forall \alpha \in a \text{ cluster device}, \exists A, B \in \text{RAR}, \text{ such that } A \preceq B \Leftrightarrow \alpha(B) \mapsto \alpha(A).$$

we say $\alpha(B)$ is true iff α can service B. Followed by defining a relation between user's requests, we can try to build a closed topology involving those requests. It will be more likely a lattice as we can ordered those requests by inclusion through the relation we just defined. As the topology inside a cluster device can be viewed as a lattice as well by considering that the intersection of subset servers in cluster, it leads to a possibility to achieve mappings during the allocating process.

A trivial approach from the order theory is to consider Birkhoff's representation theory, which says modular lattice can be represented by down-sets of its corresponding partial order composed by join-irreducible elements. Then it is easy to view the RARs as downset of the lattice of cluster topology. Then it remains to create a metric or say ranking function to compare the performance between different mapping.

4 Evaluation

All projects need to contain a serious and careful evaluation of their results. The specifics of the evaluation method (e.g., user study, experiments, formal proof review, etc.) are intrinsic to the nature of the project, so this is something that you must discuss and agree with your supervisor early in the project. Ideally, a presentation of the method and results of your evaluation should be included in its own separate section of the report.

5 Conclusions and Future work

All good projects conclude with an objective evaluation of the project's successes and failures, and suggestions for future work that can take the project further. It is important to understand that there is no such thing as a perfect project. Even the very best pieces of work have their limitations and you are expected to provide a proper critical appraisal of what you have done. Your assessors are bound to spot the limitations of your work and you are expected to be able to do the same.

Appendices

Appendices contain information that is peripheral to the main body of the report. Information typically included are things like program listings, tables, proofs, graphs, or any other material that would break up the theme of the text if it appeared in situ. Large program listings are rarely required, and should be compressed as much as possible, e.g., by printing in multiple columns and by using small font sizes, omitting inessential detail.

Literature Review

From [Sti10]:

Industrial Background: using virtual machine (VM) technology which can consolidate hardware resources to enable shared hosting platforms.

Main achievement: A formulation of the resource allocation problem in shared hosting platform for static workloads to make decisions when allocating hardware resources to service instances.

System model: Show in **Figure 1**

Resource and Service model: Each server provides several resources (e.g., CPU, RAM space, I/O bandwidth, disk space) and we consider two kinds of resource needs: *rigid* and *fluid*. A rigid need denotes that a specific fraction of a resource is required. The service cannot benefit from a larger fraction and cannot operate




Figure 1: System Model

with a smaller fraction. A fluid need specifies the maximum fraction of a resource that the service could use if alone on the server. The service cannot benefit from a larger fraction, but can operate with a smaller fraction at the cost of reduced performance. For each fluid resource need we can define the ratio between the resource fraction allocated and the maximum resource fraction potentially used. We call this ratio the *yield* of the fluid resource need. To compare yields across services with various minimum yield requirements, we define the *scaled yield* of a service as follows:

$$\text{scaled yield} = \frac{\text{yield} - \text{minimum yield}}{1 - \text{minimum yield}}.$$

After that, we can specify the assumptions and list them below:

- 1) Each service consists of only a single VM instance.
- 2) Each service has constant resource needs.
- 3) The yields of all fluid resource needs are identical.
- 4) Rigid resource need are independent from fluid resource needs
- 5) User concerns, or say higher level metrics are directly related to resource fractions allocated to services

Then we specify the aim precisely as *Maximize the Minimum Yield* under the two constraints below:

- 1) The resource capacities of the servers not be overcome.
- 2) A service inside a single VM instance should be allocated to a single server.

Finally, we recall the problem formulation as a Mixed Integer Linear Program (MILP). Say we have N services indexed by i , H servers indexed by h which each provides d types of resources. Fractions of these resources can be allocated to services. For each service i , r_{ij} denotes its resource need for resource type j ,

as a resource fraction between 0 and 1. δ_{ij} is a binary value that is 1 if r_{ij} is a rigid need, and 0 if r_{ij} is a fluid need. We use \hat{y}_i to denote the minimum yield requirement of service i , a value between 0 and 1. We define a binary variable e_{ih} that is 1 if service i runs on server h and 0 otherwise. We denote by y_{ih} the unscaled yield of service i on server h , which must be equal to 0 if the service does not run on the server. With these definitions the constraints of our linear program are as follows, with Y denoting the minimum yield:

$$\forall i, h \quad e_{ih} \in \{0, 1\}, \quad y_{ih} \in \mathbb{Q} \quad (1)$$

$$\forall i \quad \sum_h e_{ih} = 1 \quad (2)$$

$$\forall i, h \quad 0 \leq y_{ih} \leq e_{ih} \quad (3)$$

$$\forall i \quad \sum_h y_{ih} \geq \hat{y}_i \quad (4)$$

$$\forall h, j \quad \sum_i r_{ij}(y_{ih}(1 - \delta_{ij}) + e_{ih}\delta_{ij}) \leq 1 \quad (5)$$

$$\forall i \quad \sum_h y_{ih} \geq \hat{y}_i + Y(1 - \hat{y}_i) \quad (6)$$

About Vector Packing Algorithms: Resource Allocation Problem is related to the multi-dimensional version of bin-packing, or say vector packing: our service may have fluid resource needs. Attempt to place d-dimensional vectors into (at most) H bins.

From [Rai12]:

A novel approach (Wrasse solver) to resource allocation that permits the problem specification to envelope with ease. The key contribution is defining a specification language for Resource Allocation, which is expressive enough to encode a multitude of allocation problems without using any domain-specific abstractions.

Face the fact that: there may be conflicting constraints: a strategy that is optimised for performance may not necessarily meet fault-tolerance requirement.

From [Sti12]:

Focus on the theoretical model: In addition of [Sti10], as heterogeneous cases, we consider each server's each resource (or say node's dimension) has two different capacity: Elementary capacity: a single element in each dimension; Aggregate capacity: total resource capacity counting all elements.

Compared the MILP model with [Sti10]: Using different vector to represent requirement and needs separately instead of involving a binary indicator. Add one more constraint to represent the aggregate resource capacities achieved by sum the resource used from all services on this node.

Still consider the allocation under two constraints: Rigid requirement and fluid need; Both of them become ordered vector pair. Still precisely define the problem as maximize the minimum yield over all services.

From [Pan11] and [Cap01]:

Notes from Heuristics for Vector Bin Packing [Pan11] and Lower bounds and algorithms for the 2-dimensional vector packing problem[Cap01] –Focus on the new Heuristics and Experiments model, the way to set up the experiment environment.

FFD(First Fit Decreasing) Bin-centric view: Open only one bin at any time, place items into this bin from the largest suitable one. Close the bin when no item can be put in. Heuristics: Involve random choosing. Grasp[k]: pick a random one from the best k instead of the best one. Bubblesearch: the kth best is chosen with a probability proportional to $(1 - p)^k$, for a suitable p.

Analysis of FFDAvgSum, which is same as VP.CPSum in[Sti10]: For the case when some demands in a dimension always dominate the demands in the other dimension, FFDSum is more robust, the dimensions that are not scarce can be assigned smaller coefficients and have a smaller impact on the ordering. Analysis [?]: the demands across the dimensions were sampled i.i.d.

Involving classes of randomly generated instances to model the 2-dimensional case. Using two variables c,d to represent the bins capacities on different dimensions, respectively. Using W_j and V_j to represent the weights of item j on different dimensions, respectively. Using u.d.[a,b] to represent the uniformly distribution of the value for weights of item in the interval [a,b]

In [Cap01] and [Pan11], the dimensions in the first six classes are independently sampled. In classes seven and eight, the dimensions are correlated. They achieve positive correlation by setup the domain of the V_j to be a monotonic function from W_j while achieve negative correlation by setup it to be a inverse function in domain.

In [Pan11]: Trivial way to extend to multi-dimension is to set dimension (2i-1) and (2i) are correlated as dimension 1 and 2, while independent of the other dimensions. As for generate negative correlation across all dimensions: Random variables to denote the random distribution of balls in each dimension, totally 2 times of the dimensions. Multiply with an random coefficient from [10,40] and over two (Why?). Further noise by adding a random value from [0,1]. Ignored the overweight item.

From [Van09]:

This work presents an automatic virtual resource management system for Cloud infrastructures (hosting service platforms). The main advantages include: it can automate the dynamic provisioning and placement of VMs; support for heterogeneous applications and workloads; support for arbitrary application topology.

Focus on the automatic virtual resource management system:

It contains a more complicate system model, or say architecture, to simulate the resource allocation scenario which is more closed to the really world cloud infrastructures. Instead of using bins and balls to represent, this model composed by Application Environment (AE), Local Decision Module (LDM), Global Decision Module (GDM) and datacenter which contains physical machines and VMs. Briefly, we can say that AE is faced to the application associated with specific performance goals; An application-specific LDM is associated with each AE to evaluate the process baed on the current workload using service-level metric and generate a utility function of the resource allocation; A GDM is used to interact with each LDM and the real-datacenter. It is the core of this system which contains the Constraint solver to determine the management actions based on the input of LDM's utility functions and datacenter's system-level performance metrics. It works like a black box compared with LDM. We prefer to refer to the fig.2 in [Van09] of the system graph.

Two utility functions in LDM: a fixed service-level function and a dynamic resource-level function which is communicated with GDM and updated for every iteration. VM allocation vectors in LDM are used for build up the upper bound constraints given by each application.

As for the GDM: It involve two sequential process, one for determine VM allocation vectors for each application (VM Provisioning) and then for placing VMs and PMs and achieving optimisation (VM Packing). Beyond the constraints formulations based on the CPU and Memory, an interesting approach in VM Provisioning is using a coefficient to allow the administrator to trade-off between the fulfilment of the performance goals and the cost of operating the required resources. Another interesting method is illustrated during the optimisation process in VM Packing: To minimise the number of migration required to reach the new VM-to-PM assignment, or say minimise the reconfiguration to provide a strategy with few interim steps and maximum degree of parallelism.

User Guide

For projects that result in a new piece of software, you should provide a proper user guide providing easily understood instructions on how to install and use it. A particularly useful approach is to treat the user guide as a walk through of a

typical session, or set of sessions, that collectively display all the features of your software. Technical details of how the software works is rarely required. Keep it concise and simple. The extensive use of diagrams illustrating the software in action prove particularly helpful. The user guide is sometimes included as a chapter in the main body of the report, but is often better as an appendix to the main report.

References

- [Bes06] C. Bessiere. Constraint Propagation, Chapter 3 of *Handbook of Constraint Programming* (eds. Rossi, van Beek and Walsh), Elsevier, 2006.
- [Bra13] F. Brandao, J. P. Pedroso. Bin Packing and Related Problems: General Arc-flow Formulation with Graph Compression. *Technical Report DCC-2013-08*, Faculdade de Cincias da Universidade do Porto, Portugal. 2013.
- [Dav02] B. A. Davey and H. A. Priestley. *Introduction to Lattices and Order* Cambridge University Press, Cambridge, 2002.
- [Mac85] A. K. Mackworth and E. C. Freuder. The complexity of some polynomial network consistency algorithms for constraint satisfaction problems. *Artificial Intelligence*, 25:65-74, 1985.
- [Pan11] R. Panigrahy, K. Talwar, L. Uyeda, and U. Wieder. Heuristics for Vector Bin Packing *Microsoft Research*. 2011
- [Pro93] P. Prosser. Hybrid Algorithms for the Constraint Satisfaction Problem, *Computational Intelligence*, Volume 9, Issue 3, pp268-299, 1993.
- [Rai12] A. Rai, R. Bhagwan, S. Guha. Generalized Resource Allocation for the Cloud *Proceedings of the 3rd Symposium on Cloud Computing*. Oct. 2012
- [Sch04] M.P. Schellekens, The correspondence between partial metrics and semi-valuations, *Theoretical Computer Science*, Volume 315, Issue 1, 5 May 2004, 135149, 2004.
- [Sch14] M. Schellekens, T. Yu, S. Romaguera, O. Valero, *On Semivaluations and Semimodularity* Domains XI International Workshop on domain theory and applications, Paris, France, 2014.
- [Sti12] M. Stillwell, F. Vivien, H. Casanova. Virtual Machine Resource Allocation for Service Hosting on Heterogeneous Distributed Platforms *IEEE 26th International Parallel and Distributed Processing Symposium*. 2012
- [Sti10] M. Stillwell, D. Schanzenbach, F. Vivien, H. Casanova. Resource Allocation Algorithms for Virtualized Service Hosting Platforms. *Journal of Parallel and Distributed Computing*. 2010

- [Sti08] M. Stillwell, D. Schanzenbach, F. Vivien, H. Casanova. Resource Allocation using Virtual Clusters. Technical Report ICS2008-09-01, Information and Computer Sciences Dept., University of Hawai'i at Manoa, Sept. 2008